

# Runtime Organization

CMPT379

# Runtime Support

CMPT 379: Compilers

Instructor: Anoop Sarkar

[anoopsarkar.github.io/compilers-class](https://anoopsarkar.github.io/compilers-class)

# Runtime Support


- Management of runtime resources
- Correspondence between:
  - Static (compile-time) structures
  - Dynamic (run-time) structures
- Storage organization
  - Using memory to store data structures of the executing program

# Invoke the Program

- Execution of the program is initially under the control of the operating system
- When program is invoked:
  - The OS allocates space for the program
  - The code is loaded into part of the memory
  - The OS jumps to the entry point (i.e., main)

# Memory

- Compiler is responsible for:
  - Generating code
  - Orchestrating use of the data area



# Procedure Activation

- Two assumptions about programming languages
  - Execution is sequential; control moves from one point in a program to another in a well-defined order
    - Violated by concurrency
  - When a procedure is called, control always returns to the point immediately after the call
    - Violated in: Programming languages with exception

# Procedure Activation

- An invocation of procedure  $P$  is an *activation* of  $P$
- The *lifetime* of an activation of  $P$  is
  - All the steps to execute  $P$
  - Including all the steps in procedures  $P$  calls

# Procedure Activation


- The lifetime of a variable  $x$  is the portion of execution in which  $x$  is defined (until  $x$  is de-allocated)
- Note that
  - Lifetime is a dynamic (run-time) concept
  - Scope is a static concept



# Activation Trees

- Observation
  - When **P** calls **Q**, then **Q** returns before **P** returns
- Lifetimes of procedure activations are properly nested
- Activation lifetimes (sequence of function calls) can be depicted as a tree: *activation tree*

# Activation Tree



# Activation Tree

- The activation tree depends on run-time behavior
- The activation tree may be different for every program input
- Since activations are properly nested, a stack can track currently active procedures

# Stack of Active Procedures

```

int Q(int m, int n)
{
    if (n > m) {
        int i = P(m,
n);
        Q(m, i-1); Q(i
+1, n);
    }
}
    
```

Start

Q(1,3)

P(1,3)

stack

Q(1,3)
P(1,3)

Stack does not keep track of entire activation tree, just **active** procedures

# Stack of Active Procedures

```
int Q(int m, int n)
{
    if (n > m) {
        int i = P(m,
n);
        Q(m, i-1); Q(i
+1, n);
    }
}
```

Start

Q(1,3)

P(1,3)

Q(1,0)

stack


Q(1,3)
Q(1,0)

Stack does not keep track of entire activation tree, just **active** procedures

# Stack of Active Procedures

```

int Q(int m, int n)
{
    if (n > m) {
        int i = P(m,
n);
        Q(m, i-1); Q(i
+1, n);
    }
}
    
```



stack


Q(1,3)
Q(2,3)
P(2,3)

Stack does not keep track of entire activation tree, just **active** procedures

# Stack of Active Procedures

```

int Q(int m, int n)
{
    if (n > m) {
        int i = P(m,
n);
        Q(m, i-1); Q(i
+1, n);
    }
}
    
```



stack


Q(1,3)
Q(2,3)
Q(2,1)

Stack does not keep track of entire activation tree, just **active** procedures

# Stack of Active Procedures

```

int Q(int m, int n)
{
    if (n > m) {
        int i = P(m,
n);
        Q(m, i-1); Q(i
+1, n);
    }
}
    
```




stack

Q(1,3)
Q(2,3)
Q(3,3)

Stack does not keep track of entire activation tree, just **active** procedures



# Memory Organization

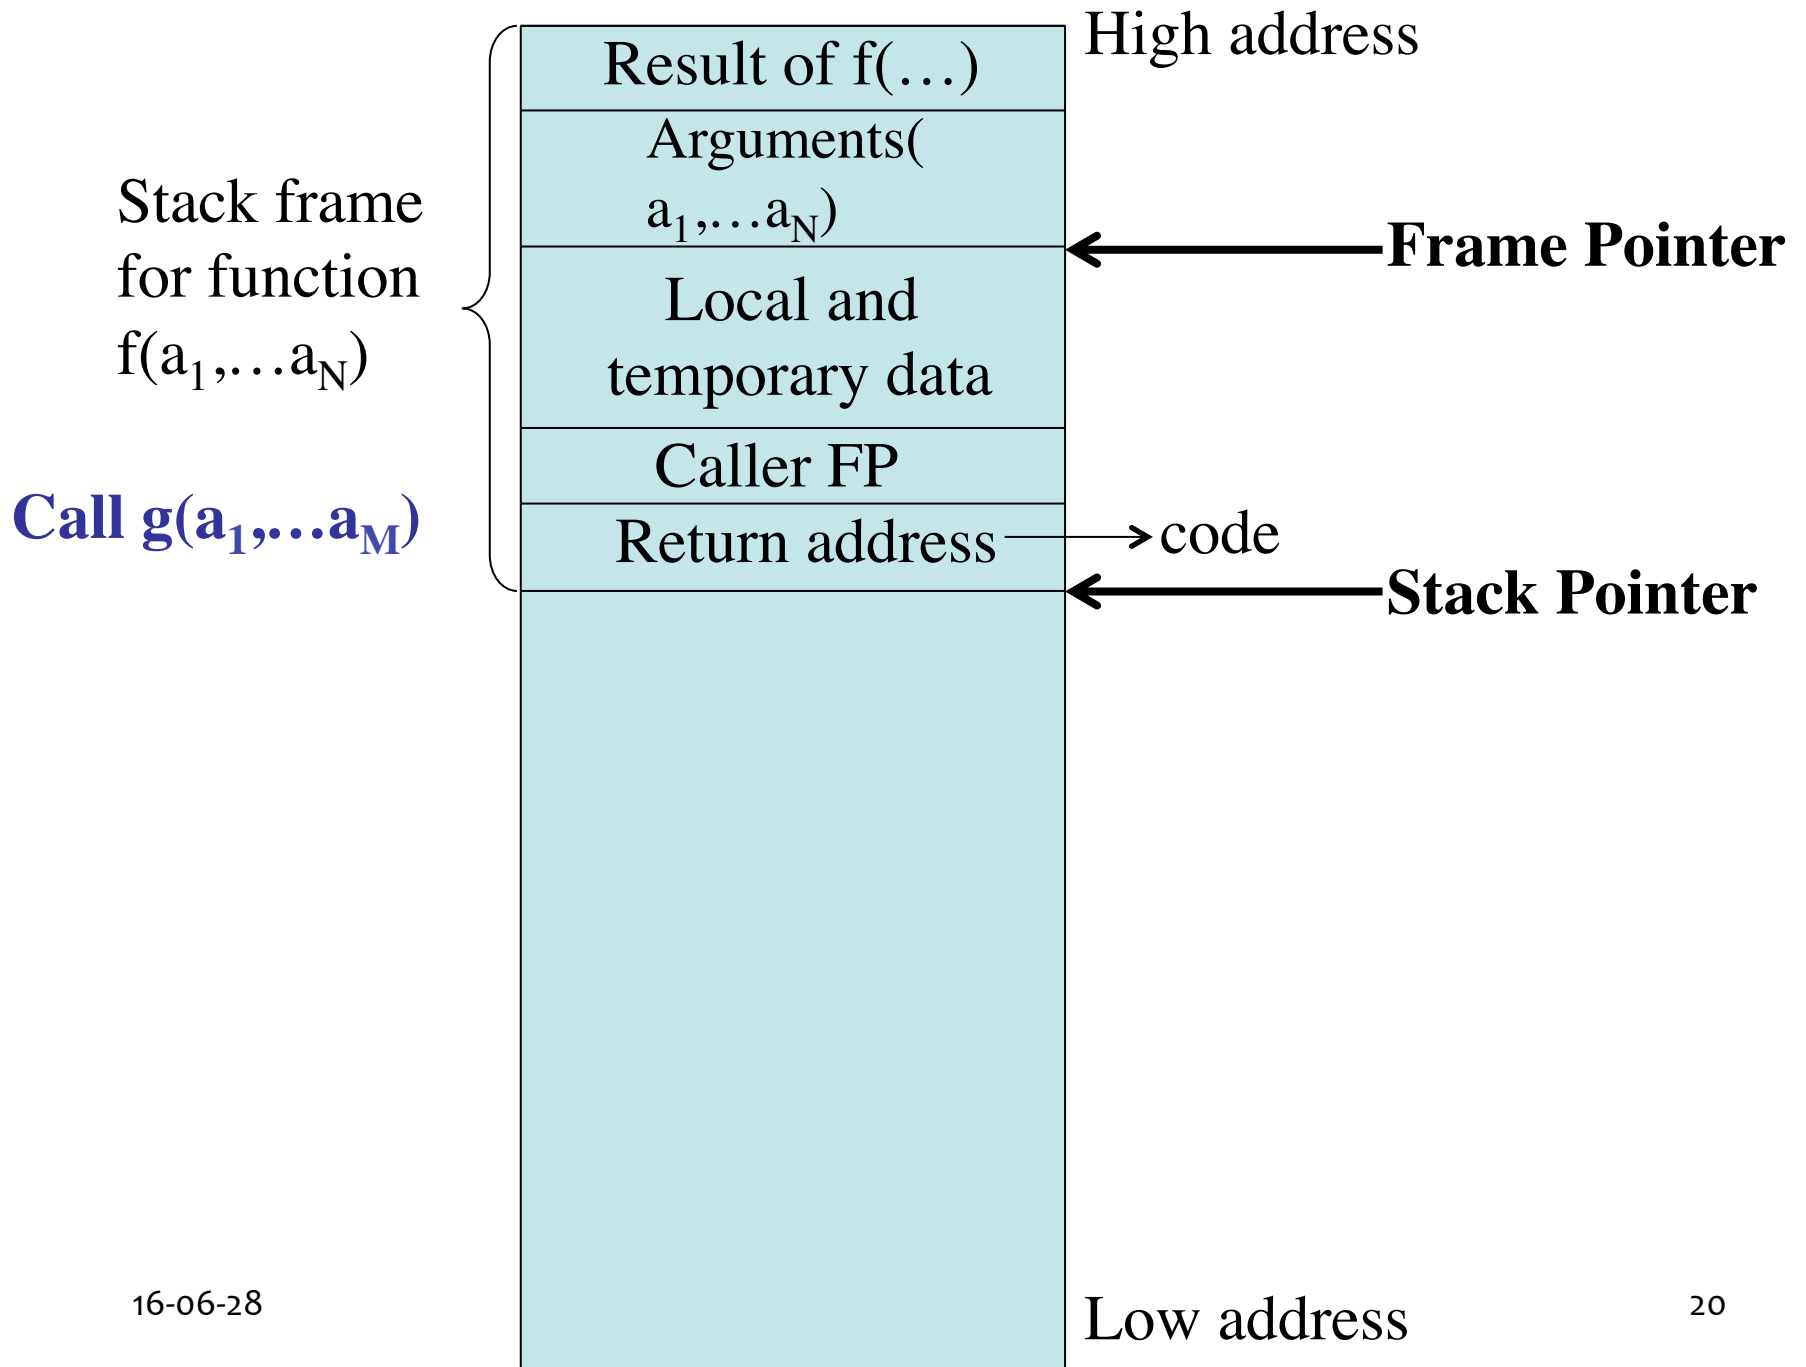


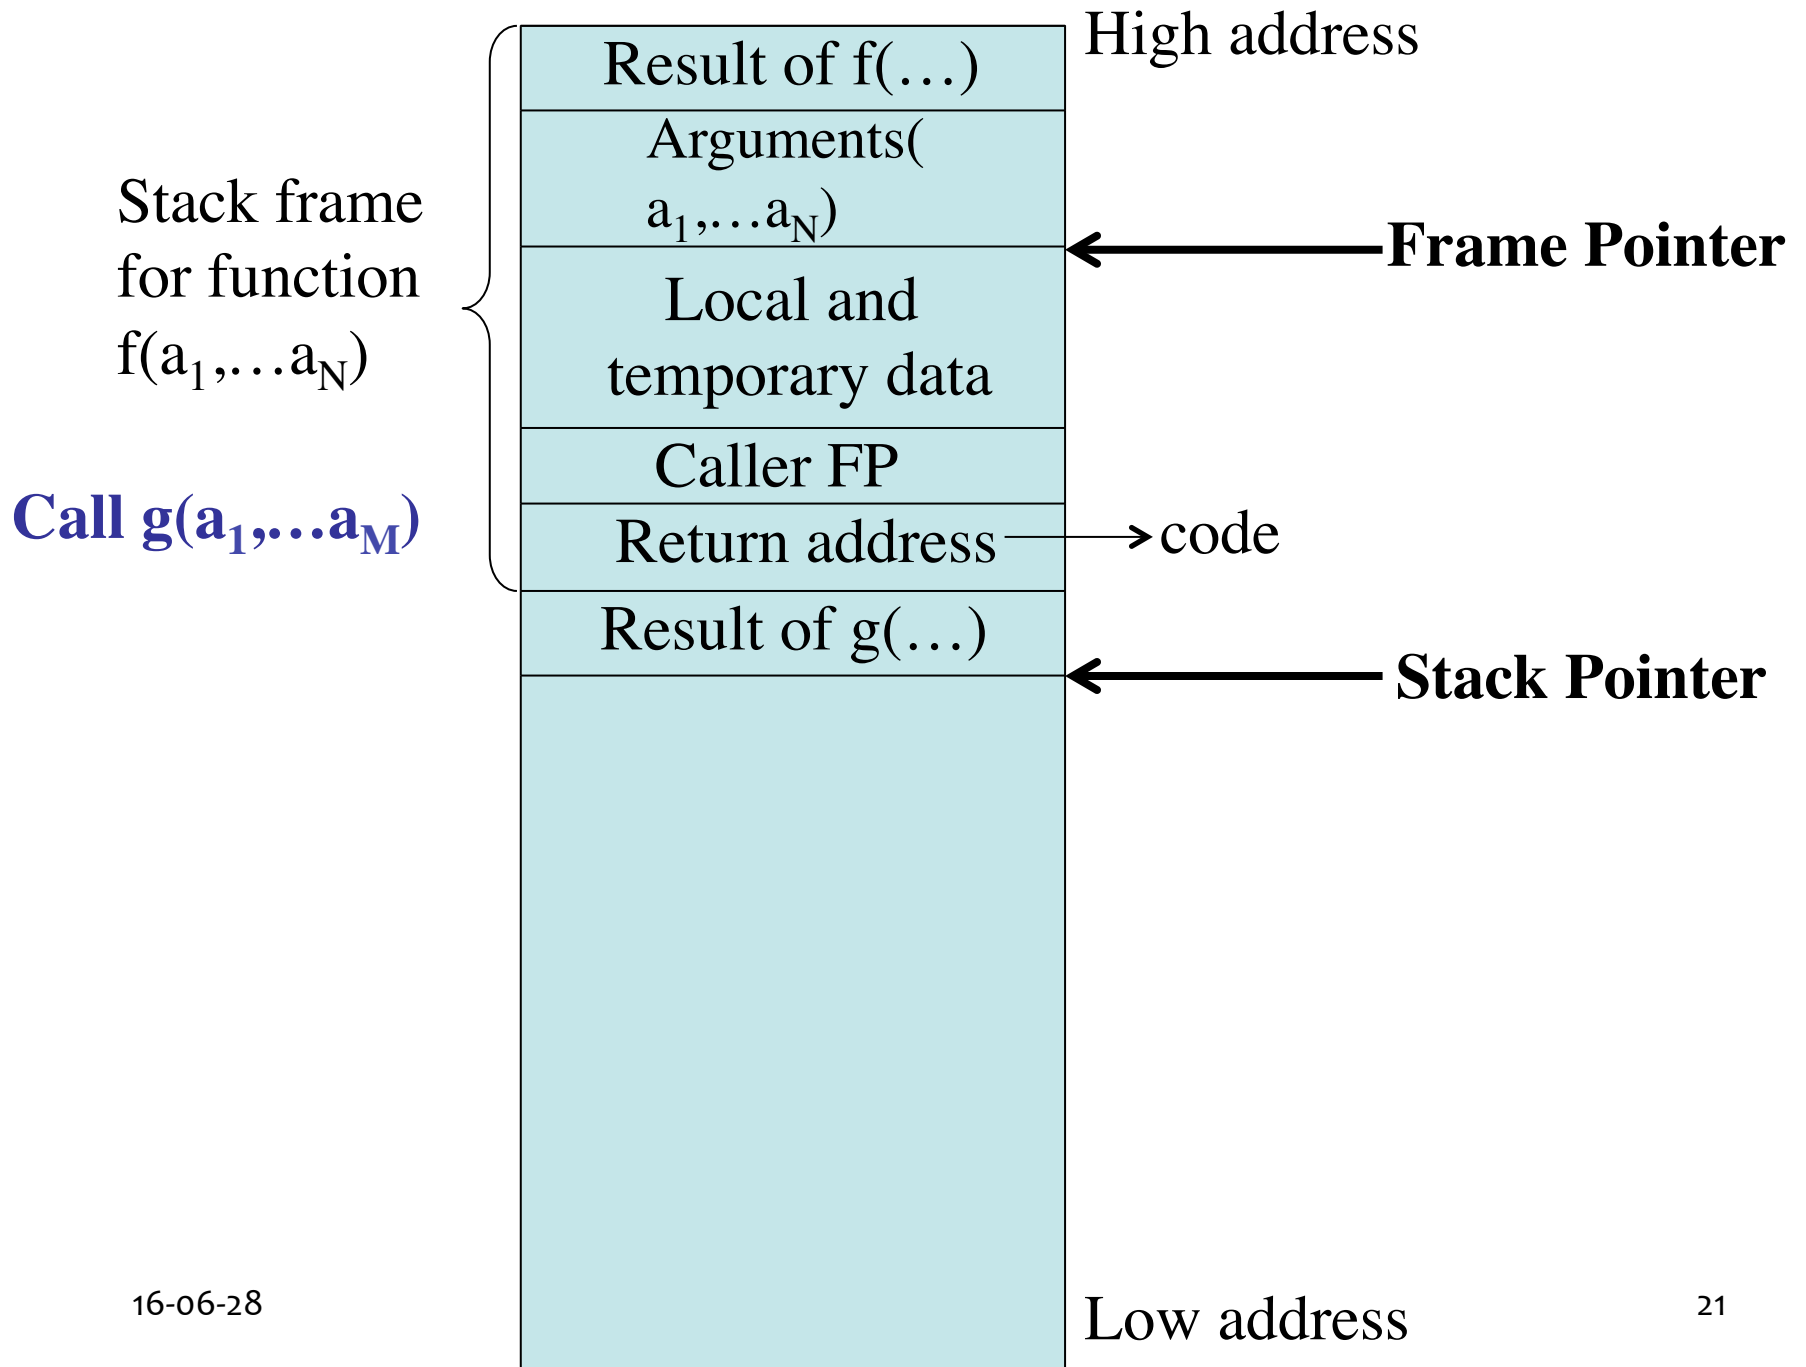
# Activation Records

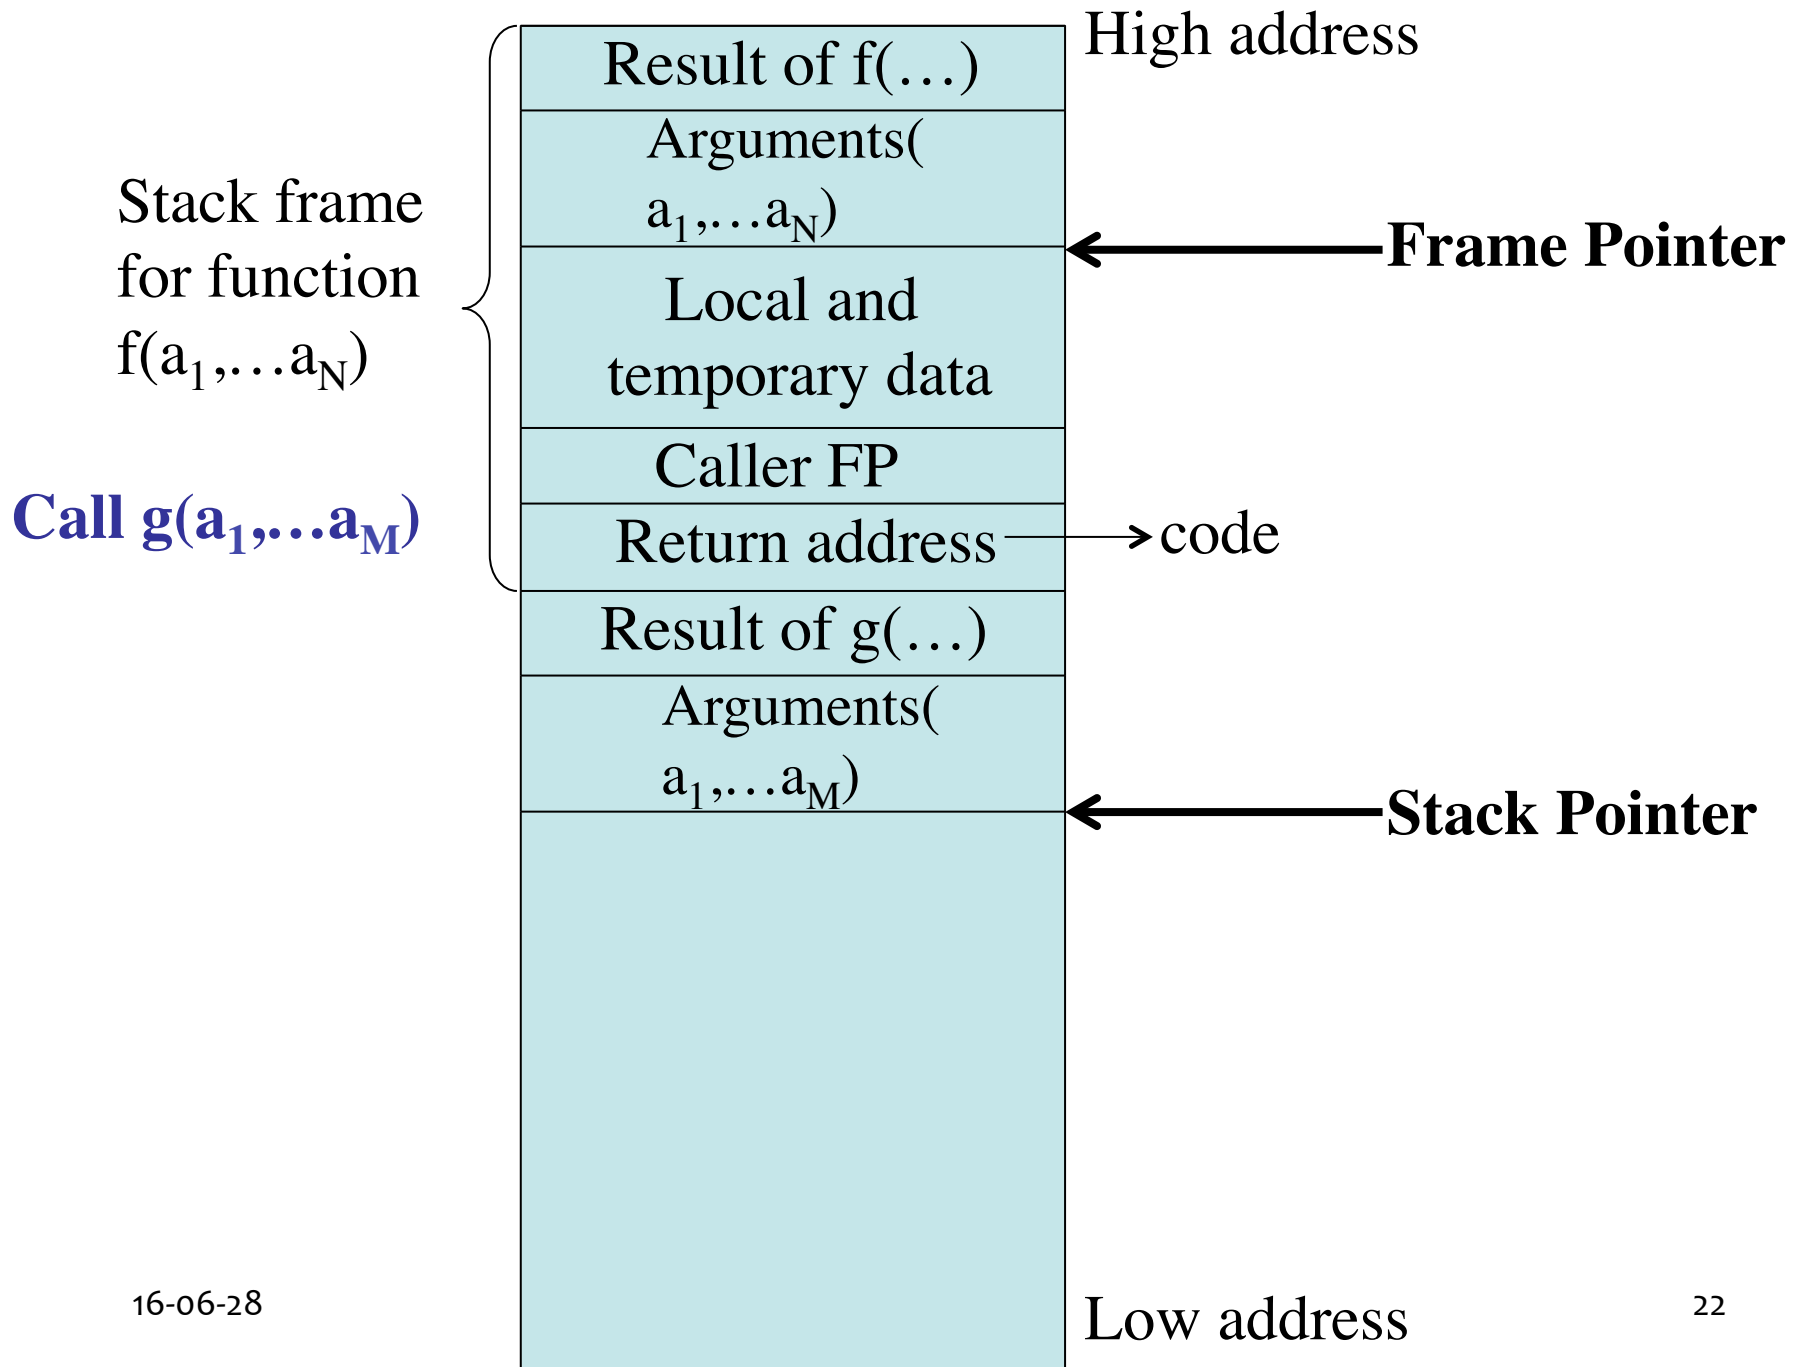
- The information needed to manage one procedure activation is called an *activation record* (AR) or *frame*
- If procedure **F** calls **G**, then **G**'s activation record contains mix of info about **F** and **G**
- **F** is suspended until **G** complete, at which point **F** resumes
- **G**'s AR contains information needed to
  - Complete execution of **G**
  - Resumes execution of **F**

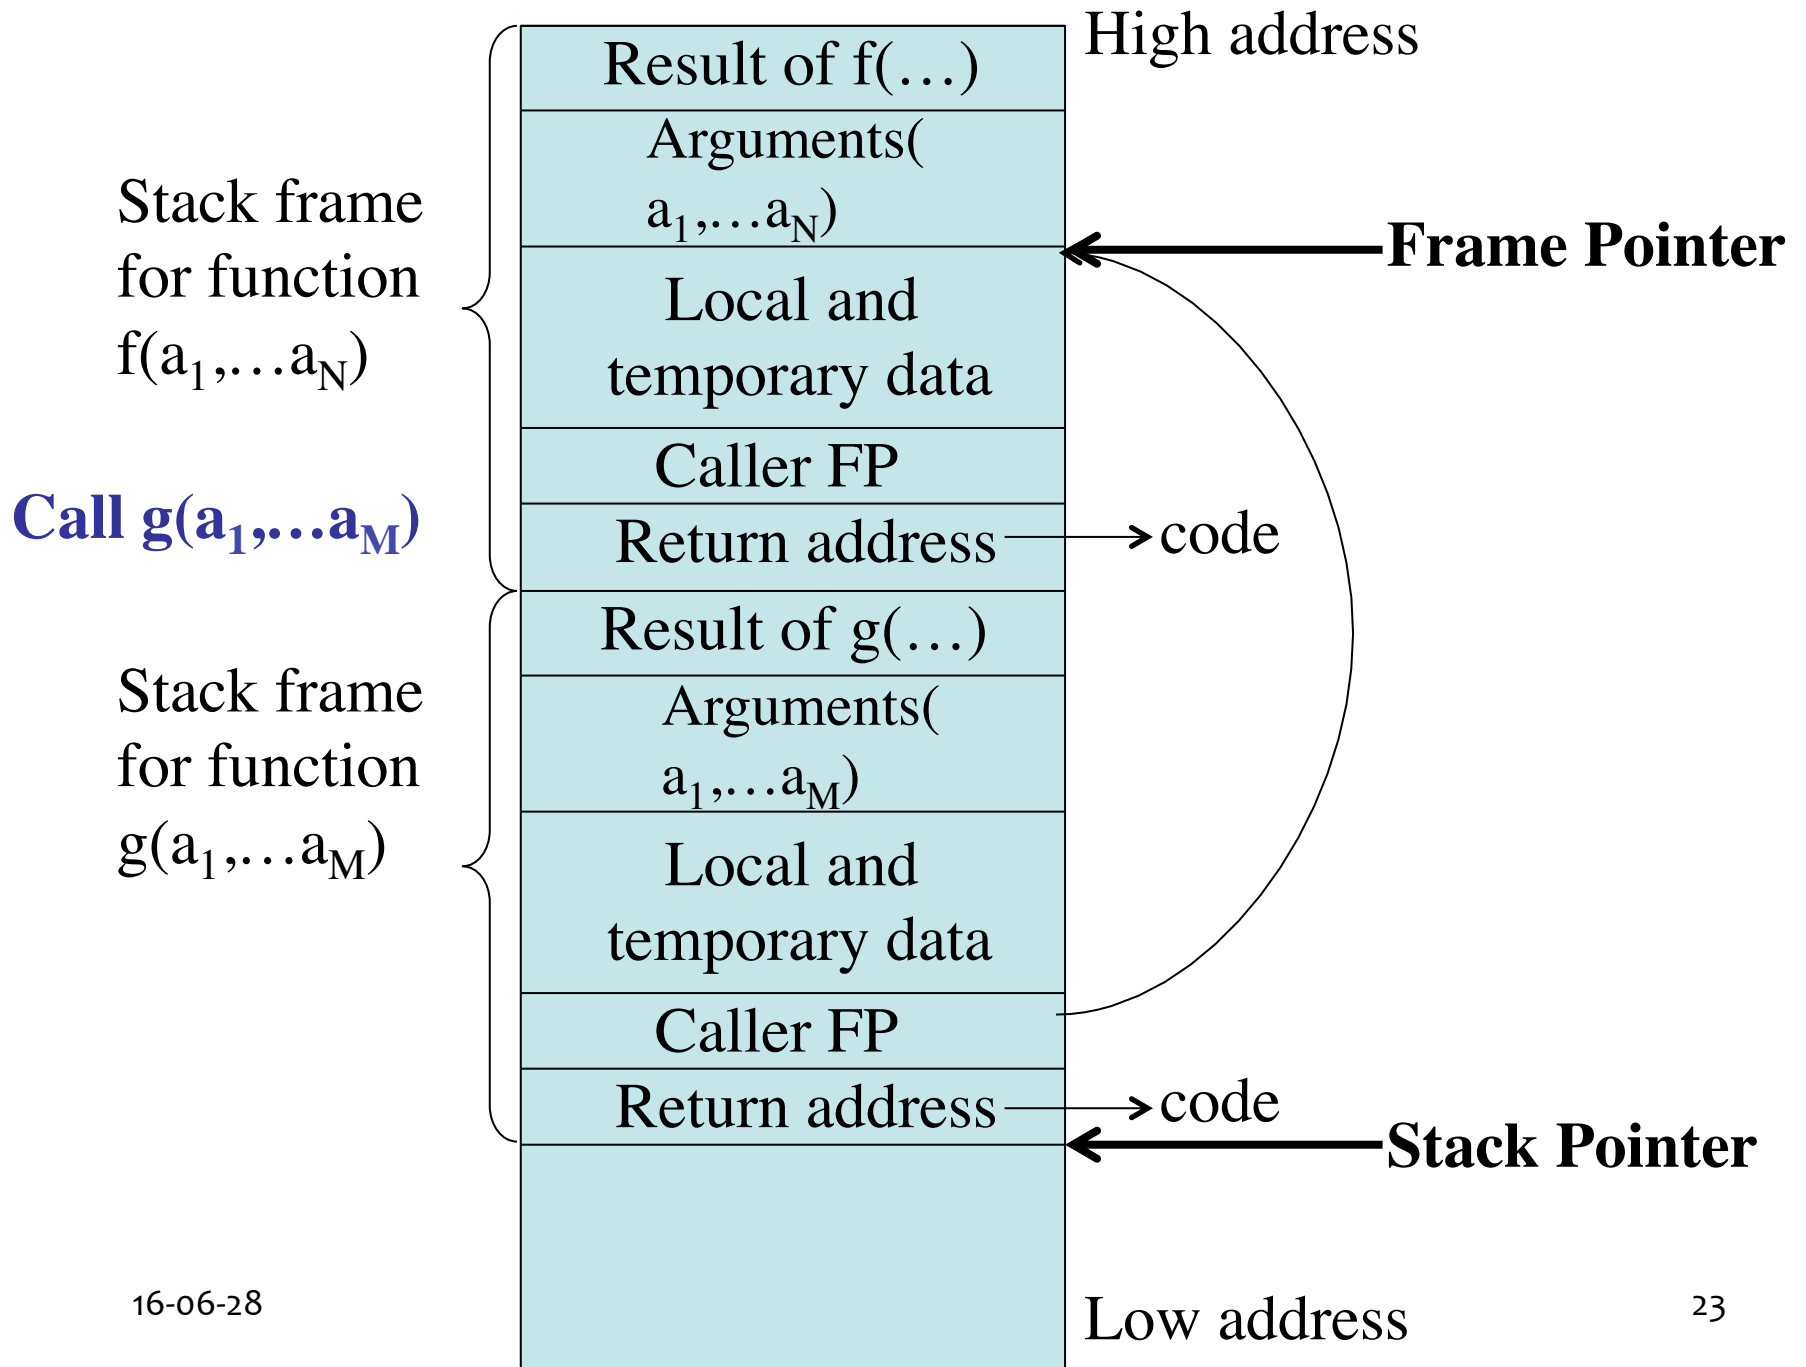
# Activation Records

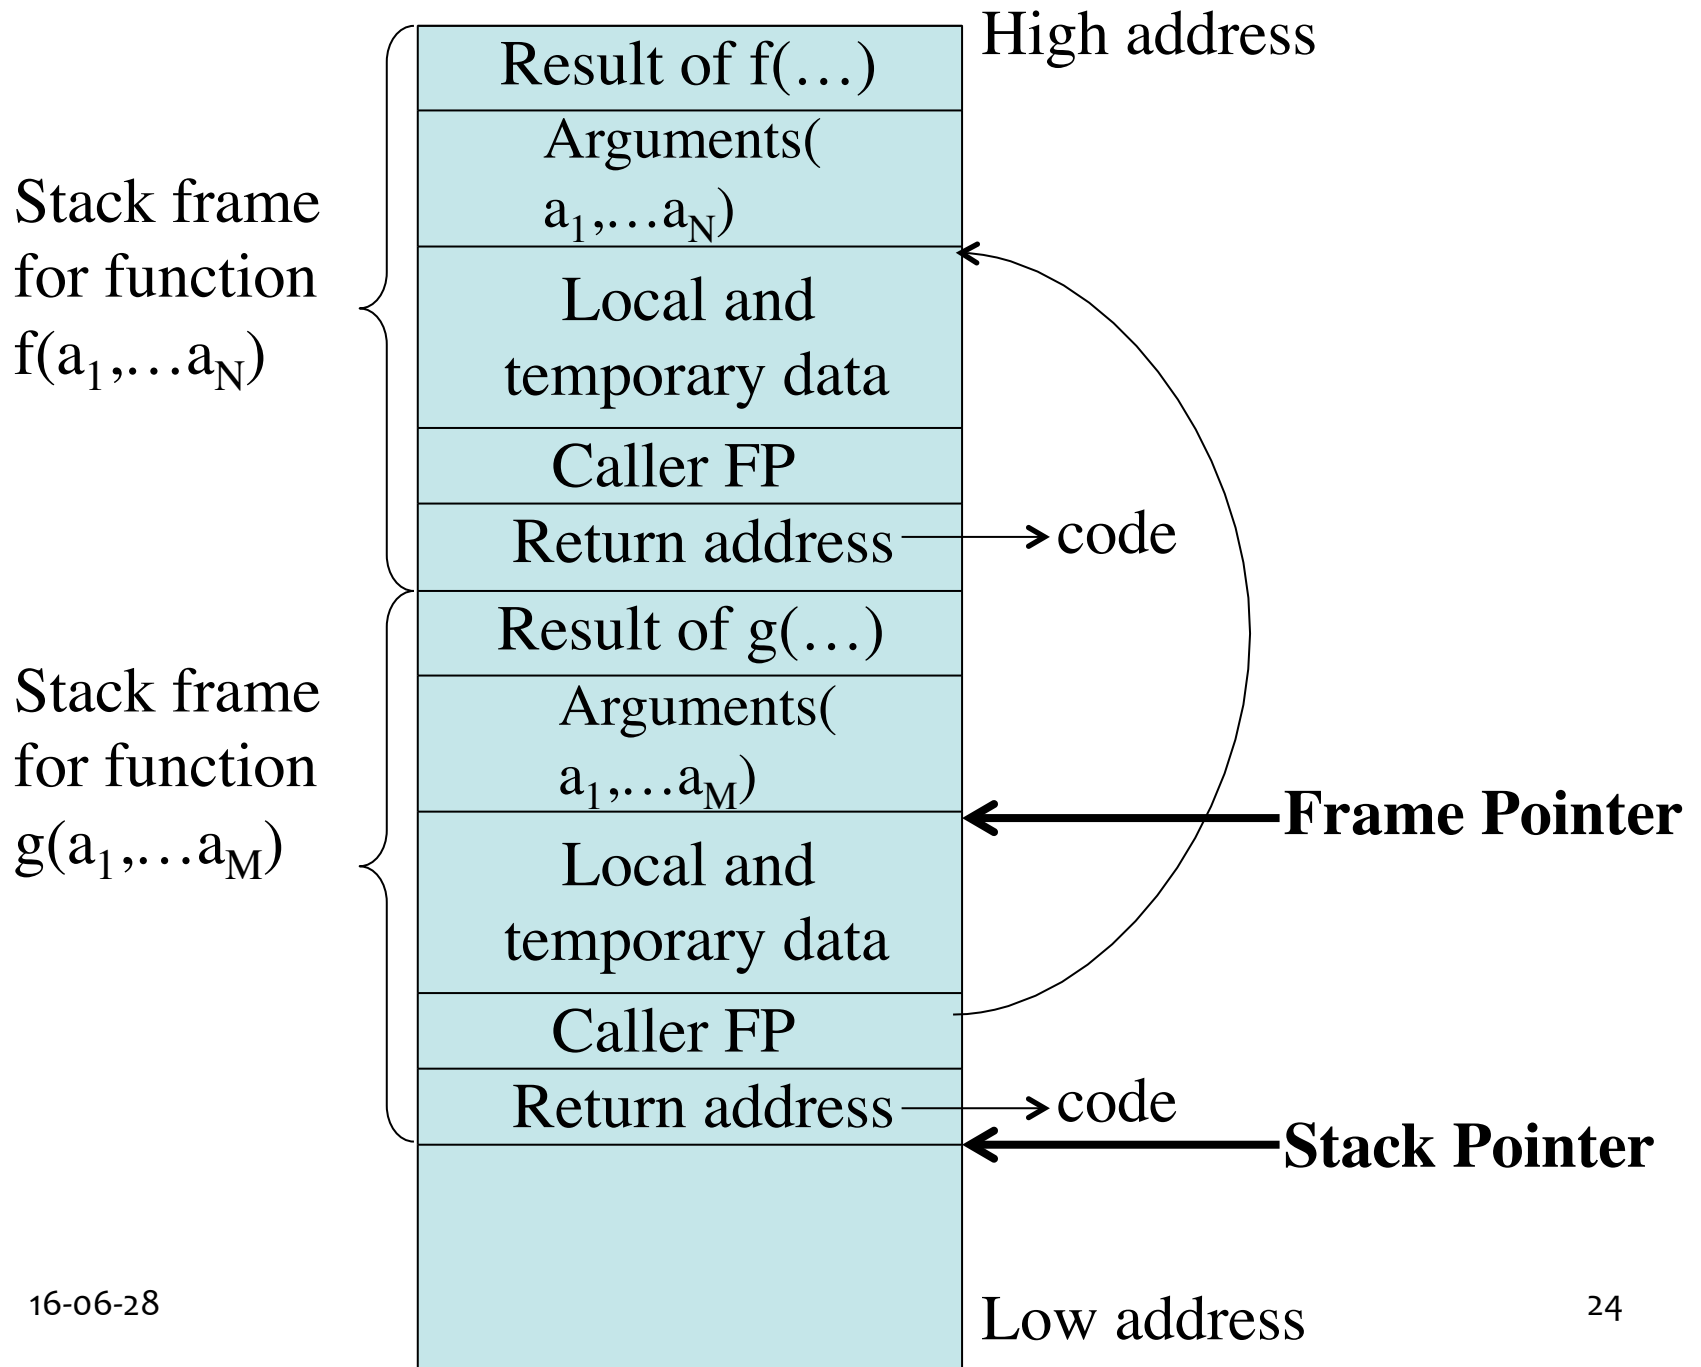
- A frame contains:
  - Control link (pointer to the caller frame)
  - Local data
  - Snapshot of machine state (important registers)
  - Return address
  - Link to global data
  - Parameters passed to function
  - Return value for the caller



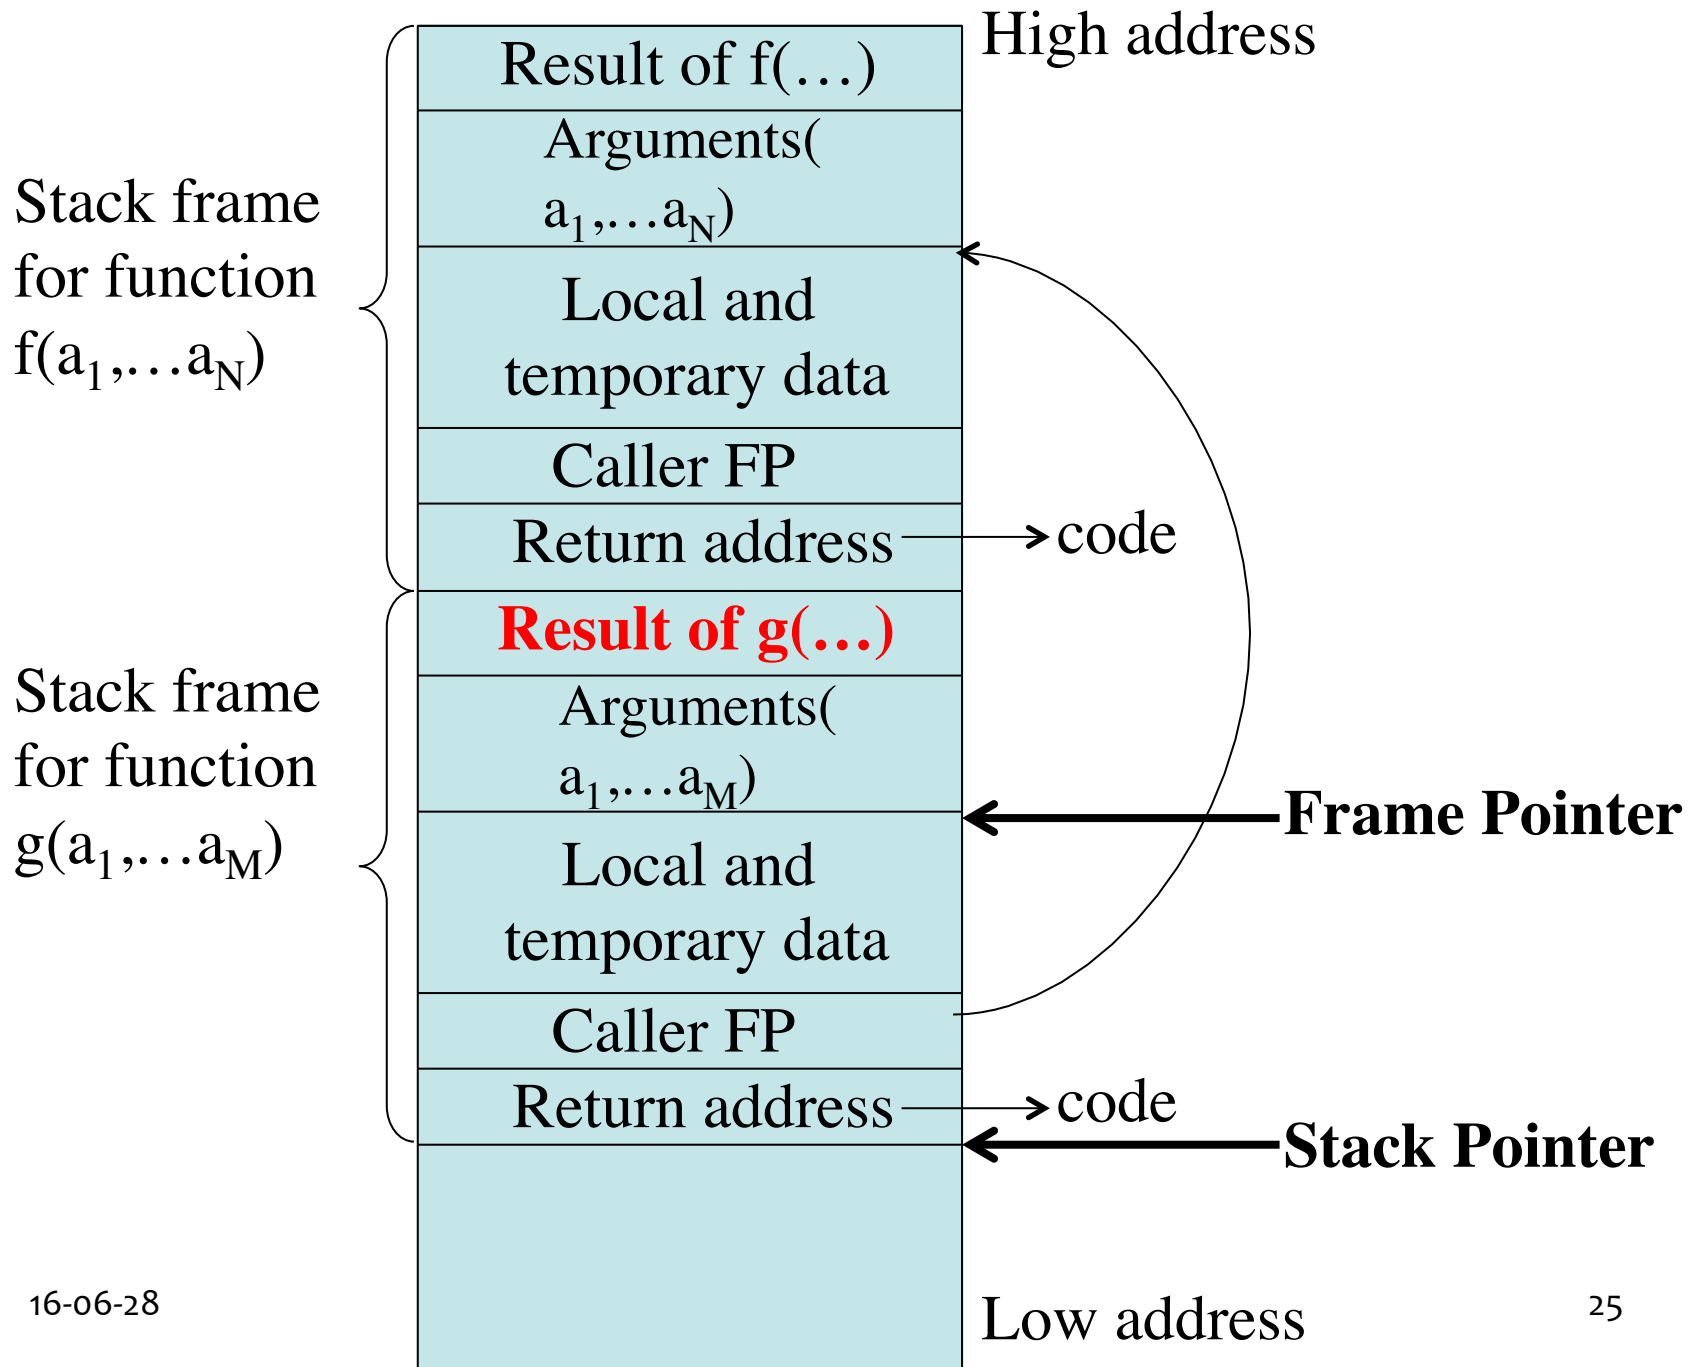


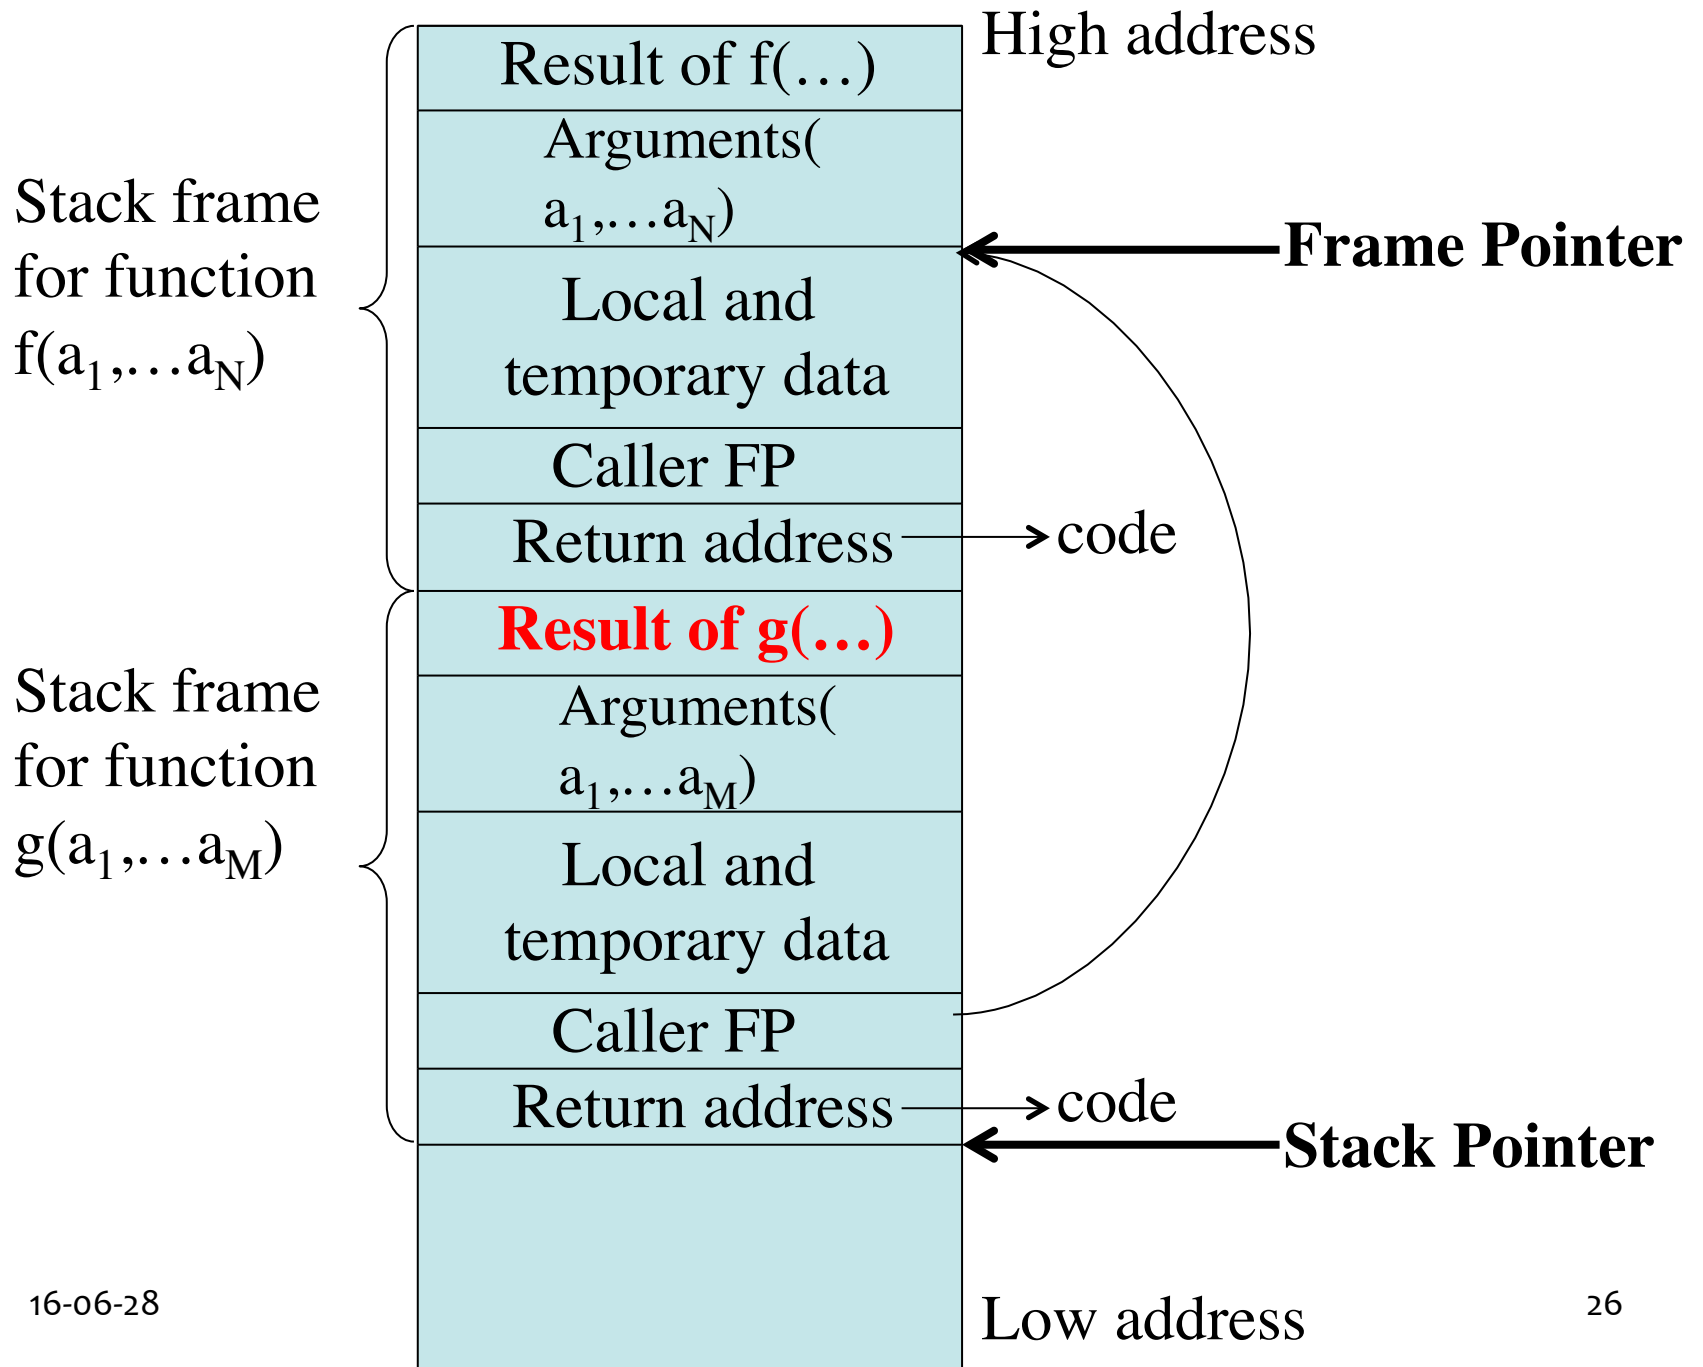


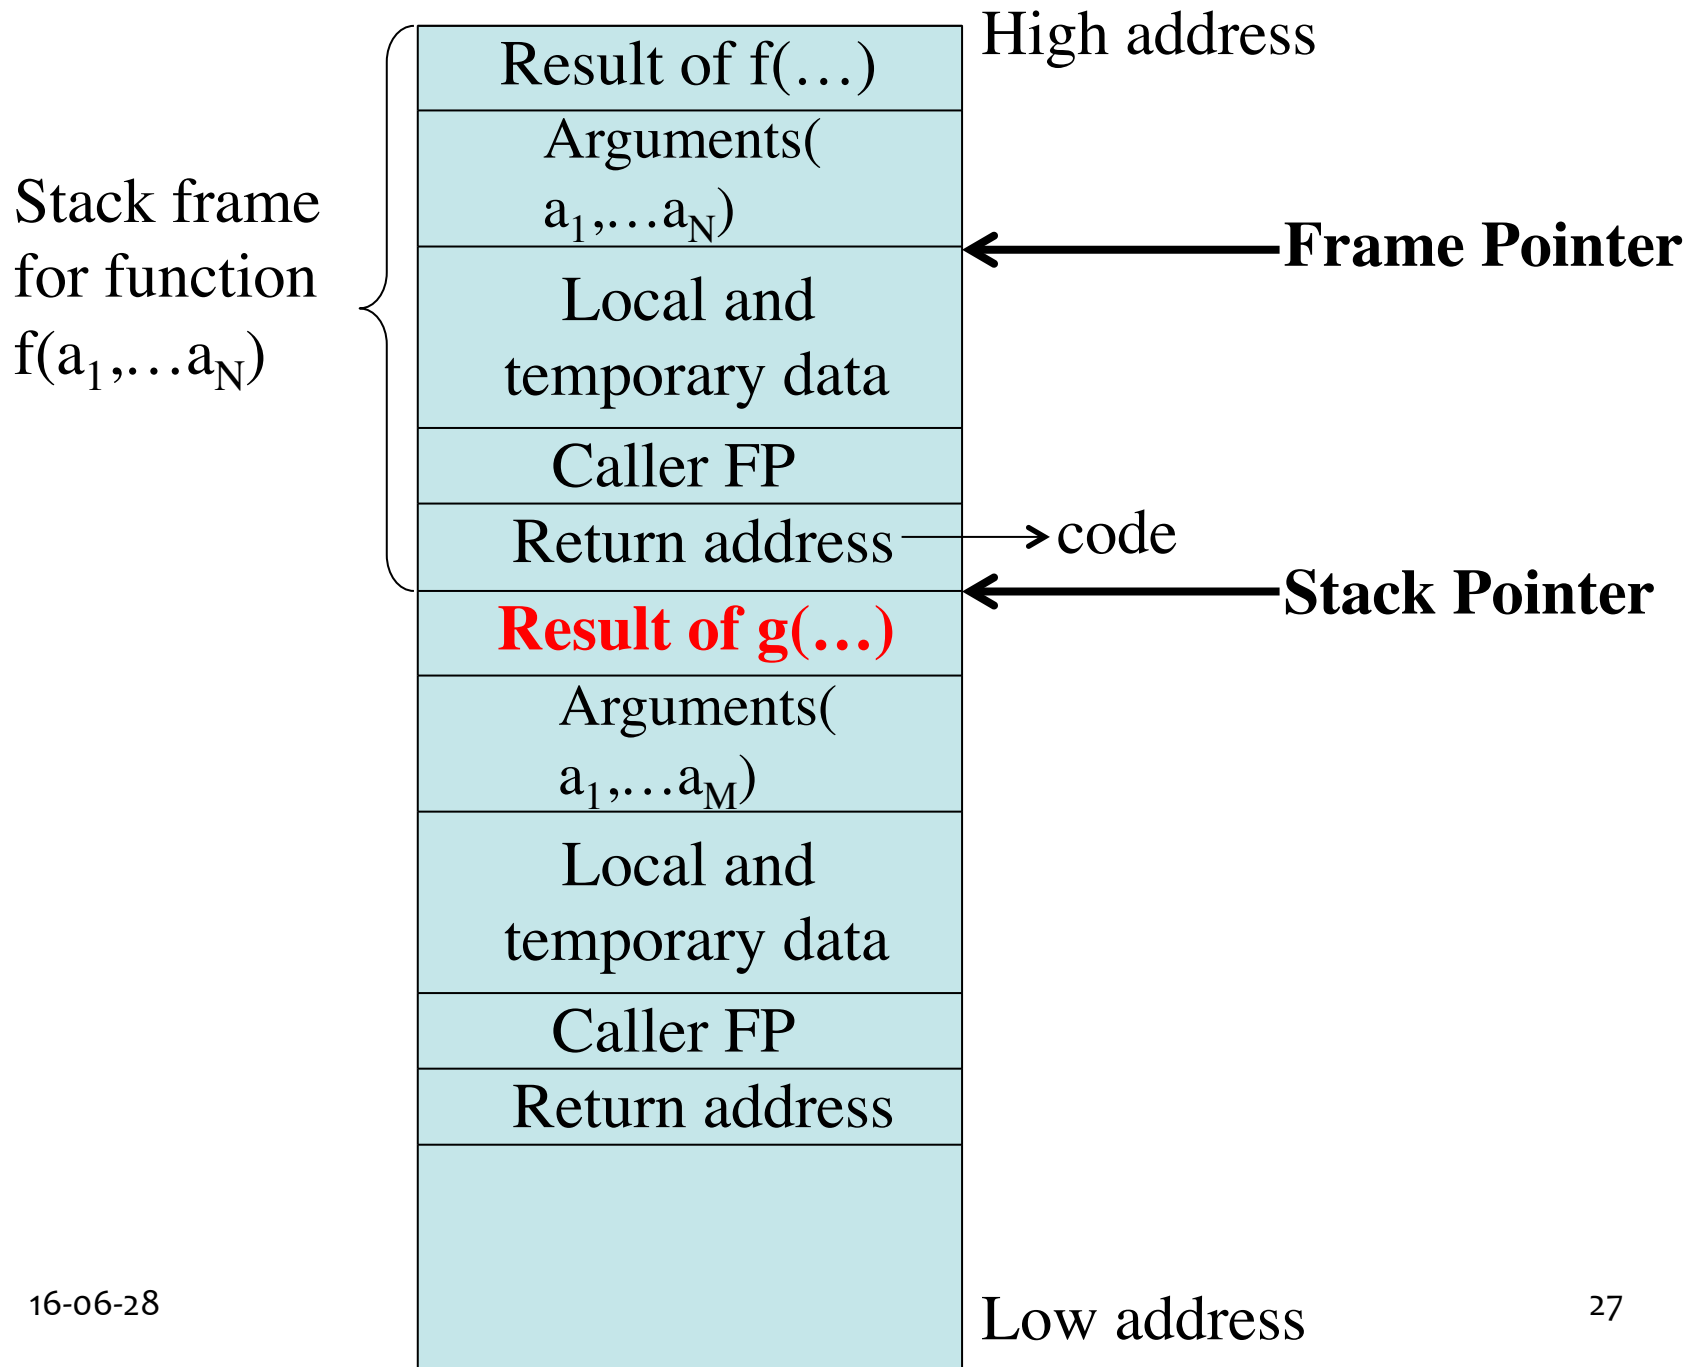








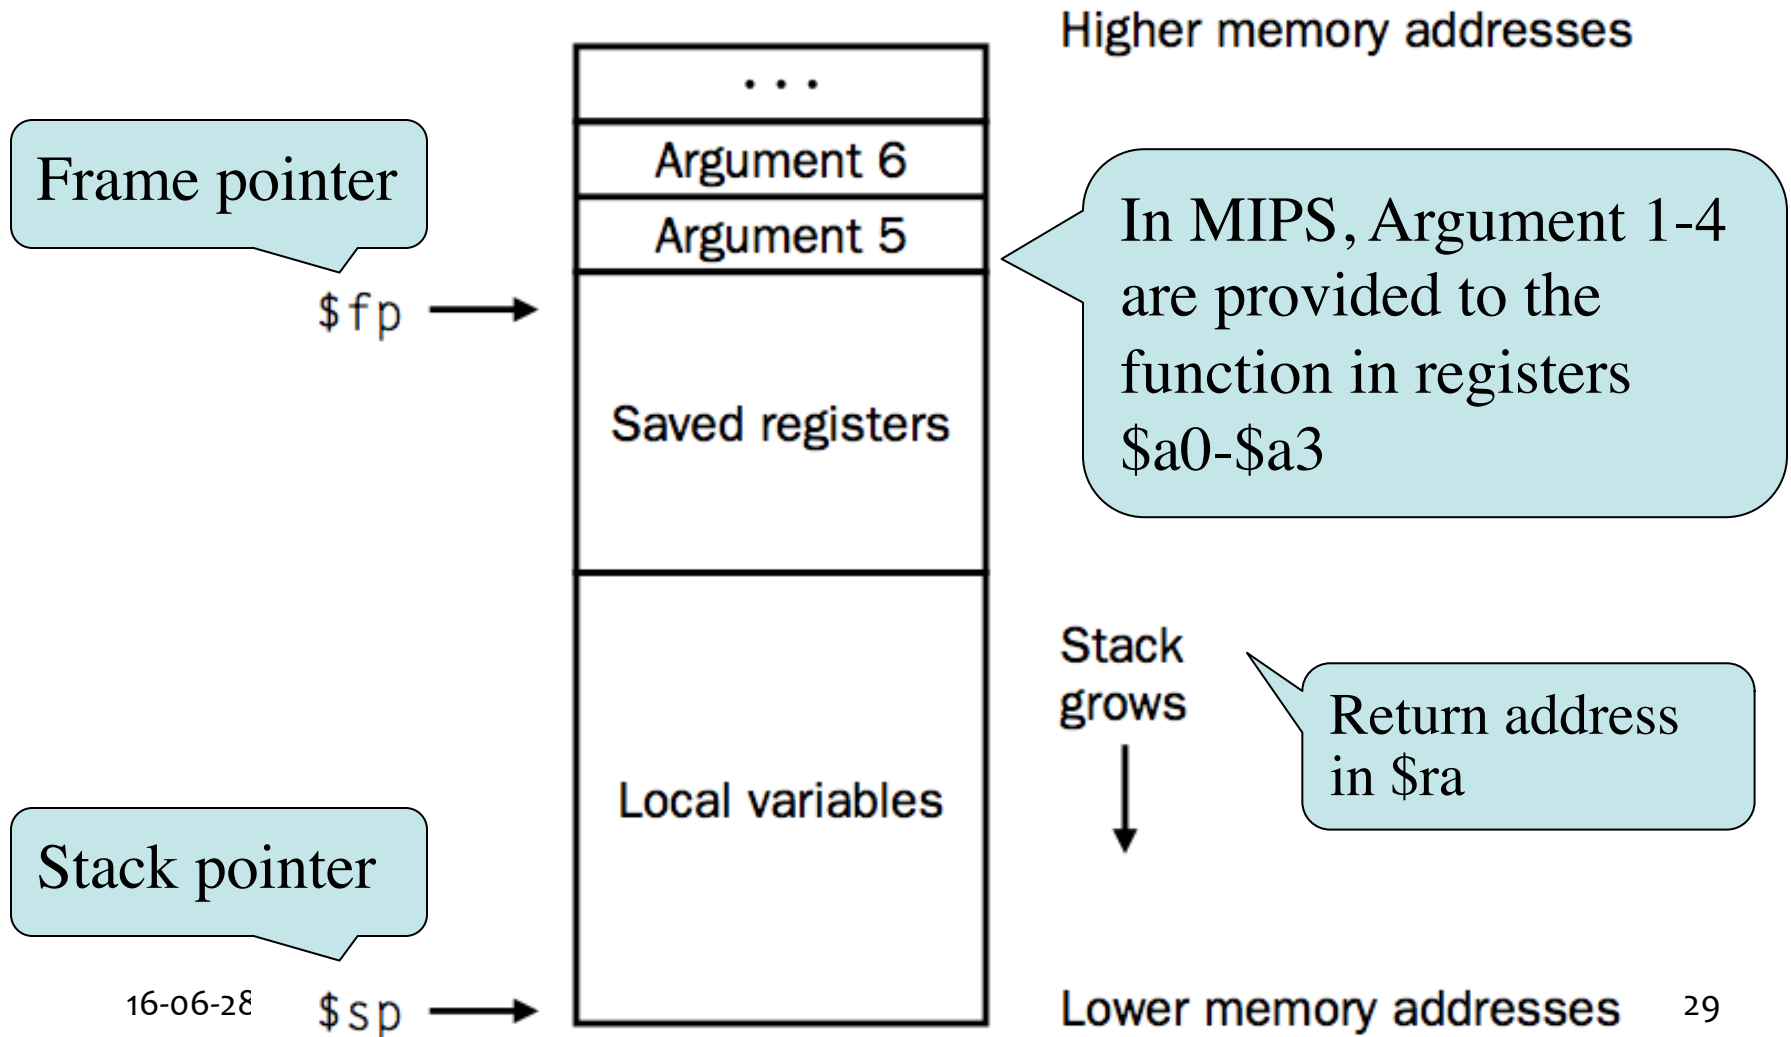




# Activation Record Organization

- There is nothing magic about this organization
  - Can rearrange order of frame elements
  - Can divide caller/callee responsibilities differently
  - An organization is better if it improves execution speed or simplifies code generation
- Real compilers hold as much of the frame as possible in registers
  - Especially the method result and arguments

# Stack frame



```
#include <stdio.h>
```

```
main ()
```

```
{
```

```
    int n = 10;
```

```
    printf("The factorial of 10 is %d\n", fact(n));
```

```
}
```

```
int fact (int n)
```

```
{
```

```
    if (n < 1)
```

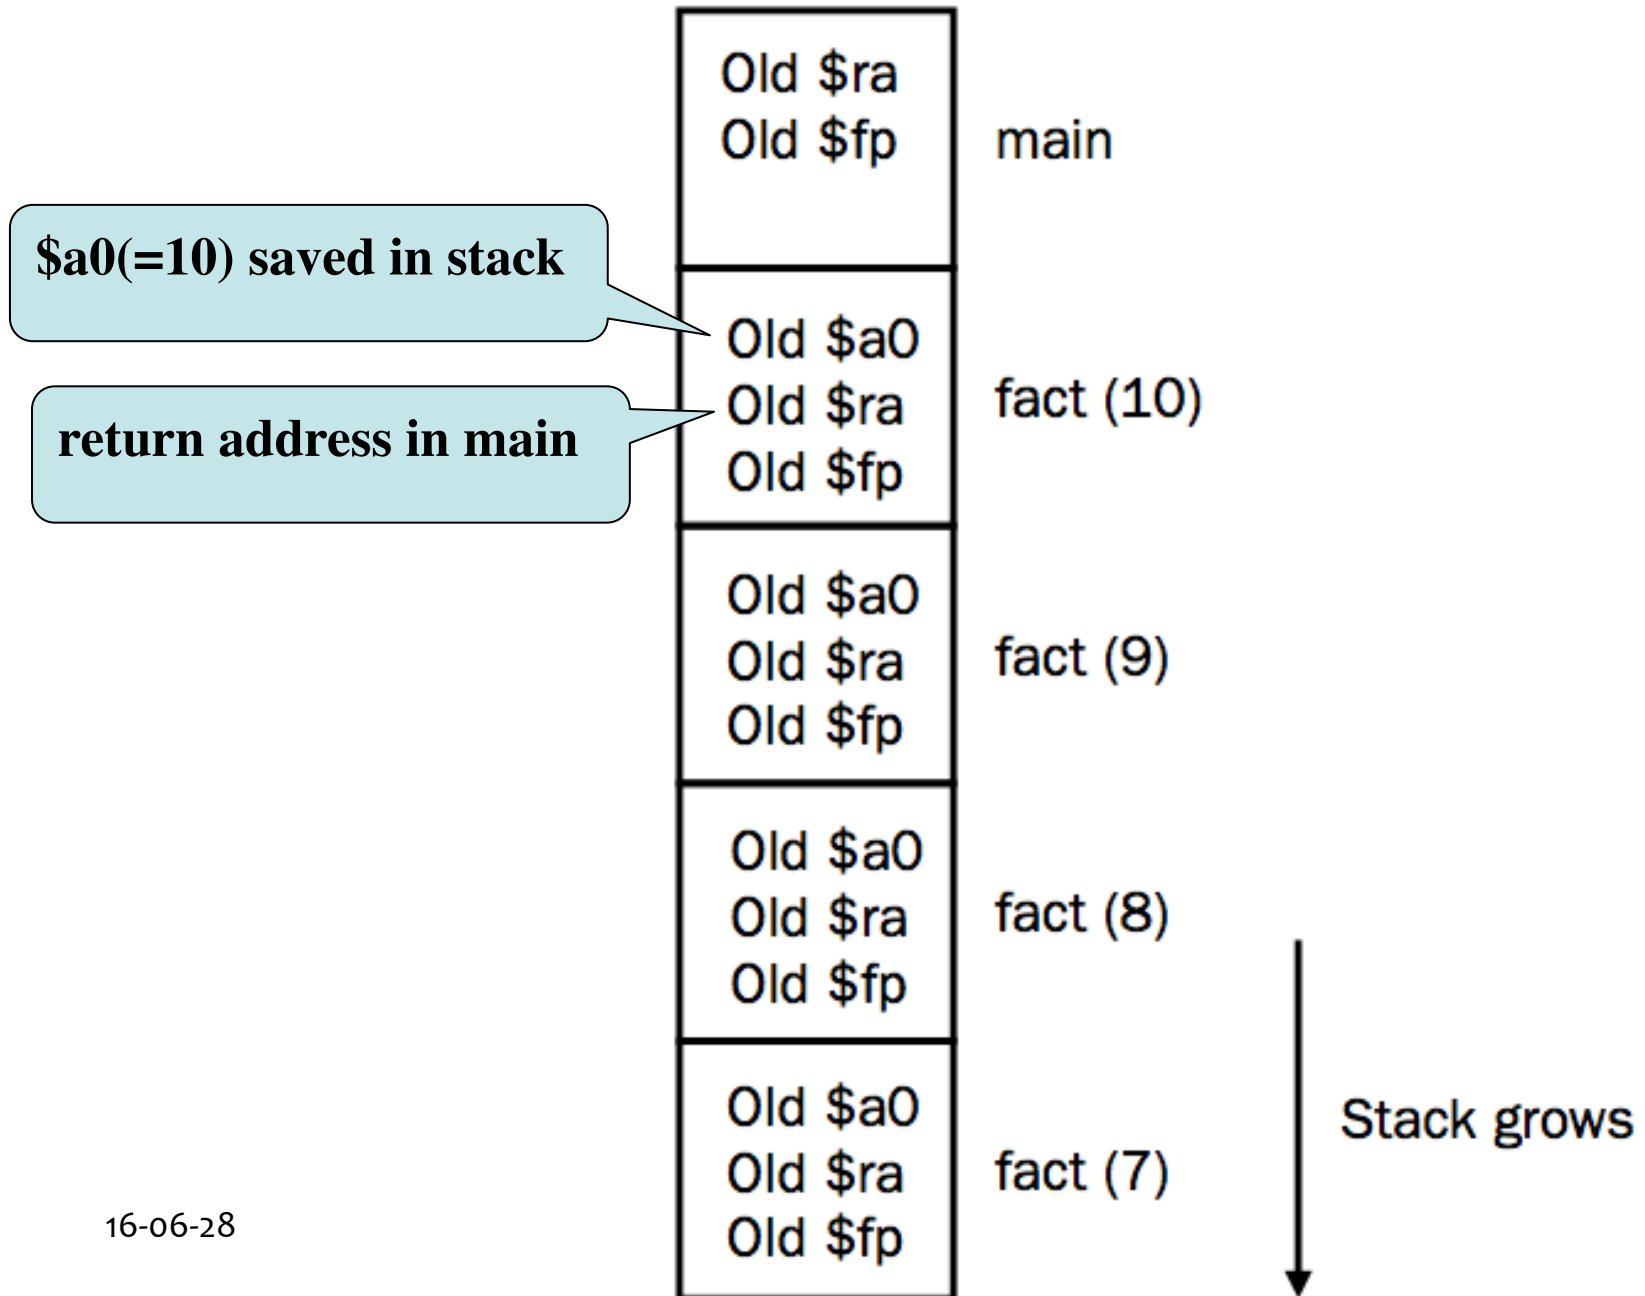
```
        return(1);
```

```
    else
```

```
        return(n * fact(n - 1));
```

```
}
```

# Stack




# Global Variables

- All references to a global variable point to the same object
  - Cannot store a global in an activation record
- Globals are assigned a fixed address once
  - Variables with fixed address are “statically allocated”
- Depending on the language, there may be other statically allocated values



# Memory Organization



# Heap Allocation

- Any value that outlives the procedure that creates it cannot be kept in AR

```
int* foo() {int * bar = new int[size]; return bar;}
```

The bar value must survive de-allocation of foo's AR

- Languages with dynamically allocated data use a heap to store dynamic data

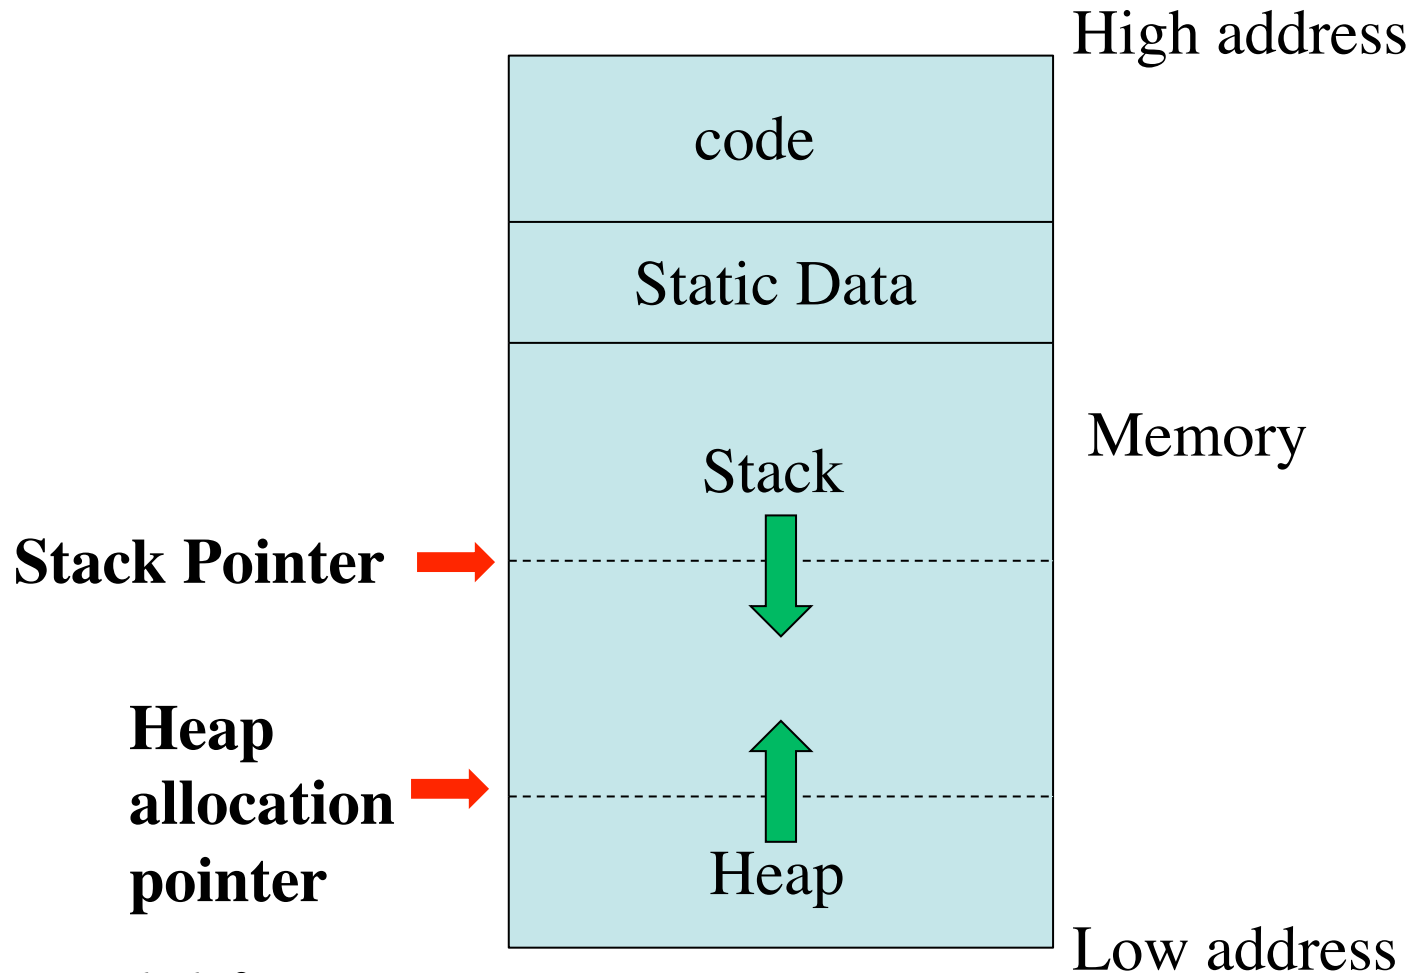
# Memory organization

- The code area contains object code
  - For many languages, fixed size and read only
- The static area contain data (not code) with fixed addresses (e.g., global data)
  - Fixed size, may be readable or writable
- The stack contains and AR for each currently active procedure
  - Each AR usually fixed size, contains locals
- Heap contains all other data
  - In C, heap is managed by *malloc* and *free*

# Heap and Stack Management

- Both the heap and stack grow
- Must take care that they do not grow into each other
- Solution: start heap and stack at opposite ends of memory and let them grow towards each other

# Memory Organization



# Alignment

- Most modern machines are 32 or 64 bit
  - 8 bits in a byte
  - 4 or 8 bytes in a word
  - Machines are either byte or word addressable
- Data is **word aligned** if it begins at a word boundary
- Most machines have some alignment restrictions
  - Or performance penalties for poor alignment

# Padding

- Example: A string

“Hello”

Takes 6 characters (including a terminating \0)

H	e	l	l	o	\0	×	×
---	---	---	---	---	----	---	---

- To word align next word, add 2 “padding” characters
- The padding is not part of the string, it’s just unused memory

# Padding

- Compilers may insert unused bytes called "padding bytes" after structure members to ensure that each member is appropriately aligned.

```
struct widget {
```

```
    char m1;
```

```
    int m2;
```

```
    char m3;
```

```
};
```

**On a word aligned machine:  
add 3 bytes of padding  
after m1 and m3**



# Summary

- Run-time support for functions
- Dealing with (potentially infinite) recursion
- Activation records for each function invocation
- Storage allocation for activation records in recursive function calls
- Stack allocation is easiest to implement while retaining recursion
- Functional PLs use heap allocation

# Extra Slides

# Storage Allocation for Functions

- Stack Allocation ✓
  - Storage for recursive functions is organized as a stack: last-in first-out (LIFO) order
  - Activation records are associated with each function activation
  - Activation records are pushed onto the stack when a call is made to the function
  - Size of activation records can be fixed or variable

# Storage Allocation for Functions

- Stack Allocation ✓
  - Sometimes a minimum size is required
  - Variable length data is handled using pointers
  - Locals are deleted after activation ends
  - Caller locals are reinstated and execution continues
  - C, Pascal and most modern programming languages

# Storage Allocation for Functions

- Heap Allocation
  - In some special cases stack allocation is not possible
  - If local variables must be retained after the activation ends
  - If called activation outlives the caller
  - Anything that violates the last-in first-out nature of stack allocation e.g. closures in Lisp and other functional PLs

# Storage Allocation for Functions

- Function Composition:  $(f \bullet g)(x) = f(g(x))$

```
class Compose {  
    fun sq (int x) { return (x * x); }  
    fun f (fun m) { return (m•h); }  
    fun h () { return sq; }  
    fun g (fun z) { return (sq•z); }  
    int main() {  
        fun v = g•h;  
        print_int((v())(3));  
    }  
}
```

# Storage Allocation for Functions

- Function Composition:  $(f \bullet g)(x) = f(g(x))$

```
class Compose {  
    fun sq (int x) { return (x * x); }  
    fun f (fun m) { return (m•h); }  
    fun h () { return sq; }  
    fun g (fun z) { return (sq•z); }  
    int main() {  
        fun v = g•h;  
        callout("print_int", (v())(3));  
    }  
}
```

$v = g \bullet h$

$v() = (g \bullet h)()$

$v() = g(h())$

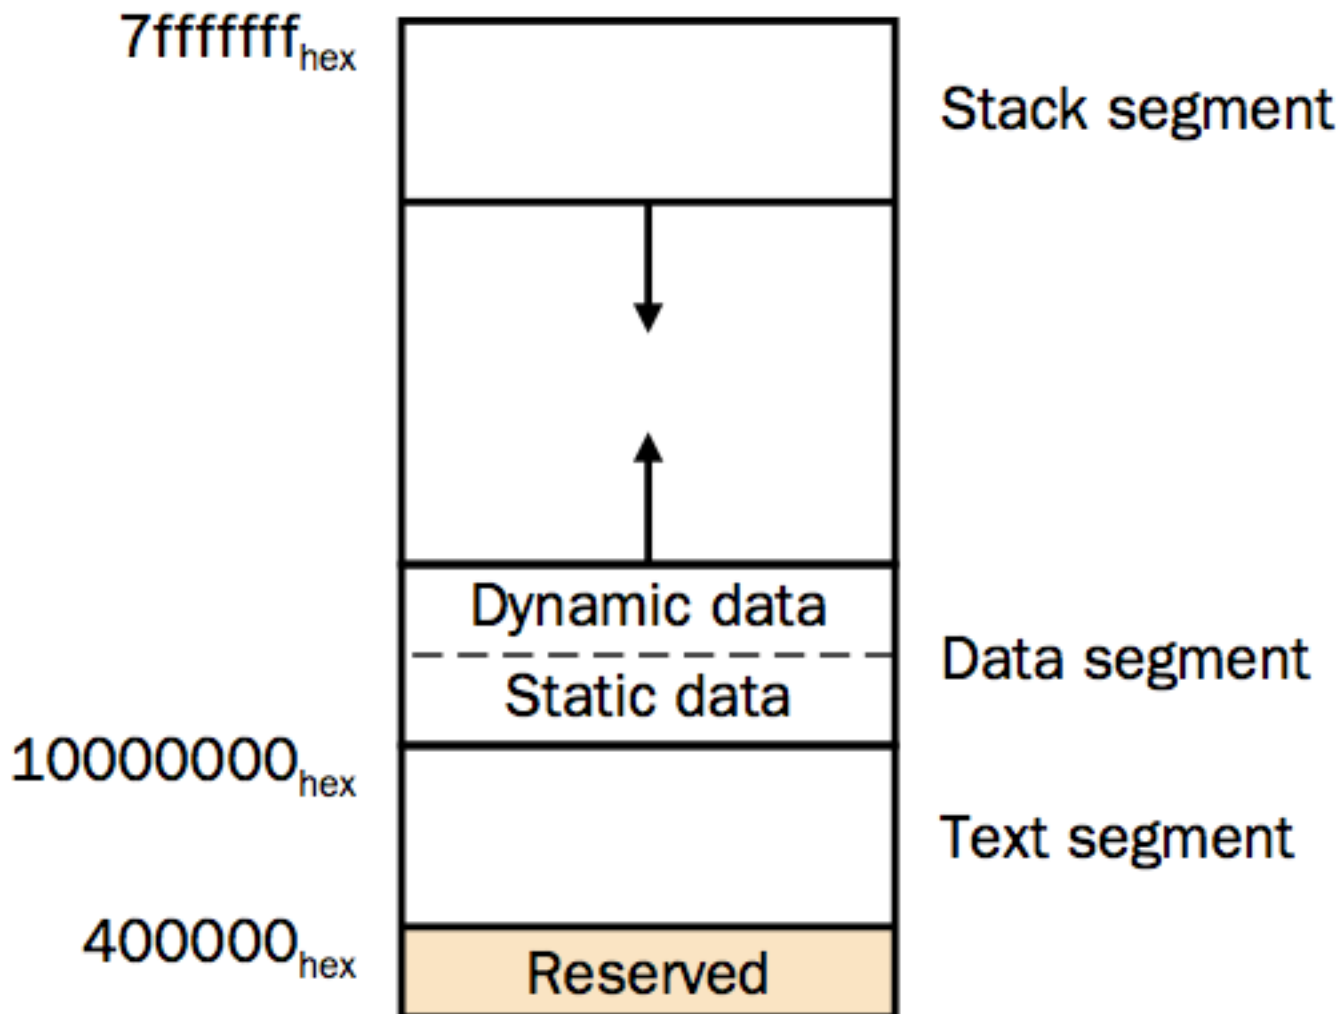
$v() = g(sq)$

$v() = (sq \bullet sq)$

$v()(3) = (sq \bullet sq)(3)$


$v()(3) = (sq(sq(3)))$

# Run-time Memory






# Stack frame



# Example: MIPS stack frame



# Parameter Passing Conventions

- Differences based on:
  - The parameter represents an r-value (the rhs of an expr)
  - An l-value
  - Or the text of the parameter itself
- Call by Value
  - Each parameter is evaluated
  - Pass the r-value to the function
  - No side-effect on the parameter

# Parameter Passing Conventions

- Call by Reference
  - Also called call by address/location
  - If the parameter is a name or expr that is an l-value then pass the l-value
  - Else create a new temporary l-value and pass that
  - Typical example: passing array elements `a[i]`

# Parameter Passing Conventions

- Copy Restore Linkage
  - Pass only r-values to the called function (but keep the l-value around for those parameters that have it)
  - When control returns back, take the r-values and copy it into the l-values for the parameters that have it
  - Fortran
- Call by Name
  - Function is treated like a macro (a #define) or in-line expansion
  - The parameters are literally re-written as passed arguments (keep caller variables distinct by renaming)

# Parameter Passing Conventions

- Lazy evaluation
  - In some languages, call-by-name is accomplished by sending a function (also called a thunk) instead of an r-value
  - When the r-value is needed the function is called with zero arguments to produce the r-value
  - This avoids the time-consuming evaluation of r-values which may or may not be used by the called function (especially when you consider short-circuit evaluation)
  - Used in lazy functional languages

# Parameter Passing Conventions

- Call-by-need
  - Similar to lazy evaluation, but more efficient
  - To avoid executing similar r-values multiple times, some languages used a memo slot to avoid repeated function evaluations
  - A function parameter is only evaluated when used inside the called function
  - When used multiple times there is no overhead due to the memo table
  - Haskell